

Descriptive Statistics 5

DataVizCheatSheet

Dr. Gordon Wright

Data Visualization Cheat Sheet

1. Central Tendency (Mean, Median, Mode)

When exploring central tendency, you can use the following charts based on the data type:

Data Type	Chart Type	Best for	Description
Nominal	Bar Chart	Mode	Displays the frequency of categories.
Ordinal	Bar Chart	Median, Mode	Bars ordered by the rank of categories.
Interval	Histogram	Mean, Median, Mode	Shows frequency distribution of values.
Interval/Ratio	Box Plot (Box-and-Whisker)	Median, Range, Outliers	Displays median, quartiles, and outliers.
Ratio	Histogram or Box Plot	Mean, Median, Mode	Shows frequency distribution or central range.

2. Measurement Types and Appropriate Charts

To visually represent different data types, the following charts are most commonly used:

Data Type	Chart Type	Best for	Description
Nominal	Pie Chart, Bar Chart	Categories, Frequencies	Compares parts to the whole (pie) or shows counts (bar).
Ordinal	Ordered Bar Chart, Dot Plot	Ranked Categories	Ordered bars or dots that reflect the ordinal nature of data.
Interval	Histogram, Line Chart	Continuous Data (e.g., temperature)	Line chart shows changes over time; histogram shows distribution.
Ratio	Scatter Plot, Histogram, Box Plot	Continuous Data with True Zero	Used for numerical data with meaningful zero.

3. Exploring Relationships Between Variables

The following charts are commonly used to explore relationships between two or more variables:

Type of Data	Chart Type	Best for	Description
Nominal vs Nominal	Stacked Bar Chart, Grouped Bar Chart	Comparing two categories	Shows distribution of one nominal variable across levels of another.
Ordinal vs Ordinal	Ordered Bar Chart, Stacked Bar Chart	Comparing two ranked categories	Displays trends in ordered categories.
Nominal/Ordinal vs Interval/Ratio	Box Plot, Violin Plot	Comparing groups for continuous variables	Box plots show central tendency and spread; violin plots show distribution.
Interval vs Interval	Scatter Plot, Line Chart	Correlation or relationship between two continuous variables	Scatter plots show correlation or pattern; line charts show trend over time.

Type of Data	Chart Type	Best for	Description
Ratio vs Ratio	Scatter Plot, Bubble Chart	Strength and direction of relationships	Scatter plots can show strength, direction, and outliers; bubble charts add size as a variable.

4. Trends Over Time

If you are looking at trends over time for different data types, consider the following:

Data Type	Chart Type	Best for	Description
Interval/Ratio	Line Chart	Continuous time-based data	Tracks changes over time (e.g., temperature, stock prices).
Nominal	Stacked Line Chart, Stacked Bar Chart	Time-based categories	Shows changes in category proportions over time.

5. Comparing Distributions

When comparing the distribution of data between groups or categories, the following charts are useful:

Data Type	Chart Type	Best for	Description
Interval/Ratio	Histogram, Box Plot	Distribution of continuous data across groups	Histograms show frequency; box plots show median and spread.
Nominal/Ordinal	Bar Chart, Pie Chart	Comparing category proportions	Bar and pie charts compare categories' proportions.

6. Correlation and Regression

To explore the correlation between two continuous variables or to visualize a regression model:

Type of Data	Chart Type	Best for	Description
Interval/Ratio	Scatter Plot with Regression Line	Correlation or regression analysis	Scatter plots visualize relationships; a regression line adds a trend.
Nominal/Ordinal vs Interval/Ratio	Faceted Scatter Plot, Box Plot	Comparing groups with continuous variables	Shows how groups differ based on continuous measures.

Self-Assessment Questions

1. What type of chart would you use to display the frequency of different job titles (e.g., teacher, engineer, doctor)?
 - ANSWER HERE
2. How would you visualize the relationship between a person's age and their income?
 - ANSWER HERE
3. If you were comparing the median salary of employees across different education levels (e.g., high school, bachelor's, master's), which chart type would you use?
 - ANSWER HERE
4. Which chart would best show the distribution of exam scores in a class?
 - ANSWER HERE
5. What chart would you use to see how sales numbers have changed over the last 10 years?
 - ANSWER HERE

This 'cheat' sheet should help you decide on the most appropriate charts for exploring central tendency, different measurement types, and relationships in their data. It also includes self-assessment questions to ensure you get a good grip of these concepts.

Histogram

Definition

A histogram is an accurate graphical representation of the distribution of a numeric variable. It takes as input numeric variables only. The variable is cut into several bins, and the number of observation per bin is represented by the height of the bar.

Here is an example showing the distribution of the night price of Rbnb apartments in the south of France. Price range is divided per 10 euros interval. For example, there are slightly less than 750 appartements with a night price between 100 and 110 euros:


```
# Libraries
library(ggplot2)

# Load dataset from github using base R
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset_1.csv")

# Ensure the data is in a proper data frame and filter
filtered_data <- subset(data, price < 300)

# Create the histogram using ggplot2
ggplot(filtered_data, aes(x = price)) +
  geom_histogram(breaks = seq(0, 300, 10), fill = "#69b3a2", color = "#e9ecef", alpha = 0.9) +
  ggtitle("Night price distribution of Airbnb apartments") +
  theme_minimal()
```

Night price distribution of Airbnb apartments



What for

Histograms are used to study the distribution of one or a few variables. Checking the distribution of your variables one by one is probably the first task you should do when you get a new dataset. It delivers a good quantity of information. Several distribution shapes exist, here is an illustration of the 6 most common ones:

Examples

```
# Libraries
library(ggplot2)

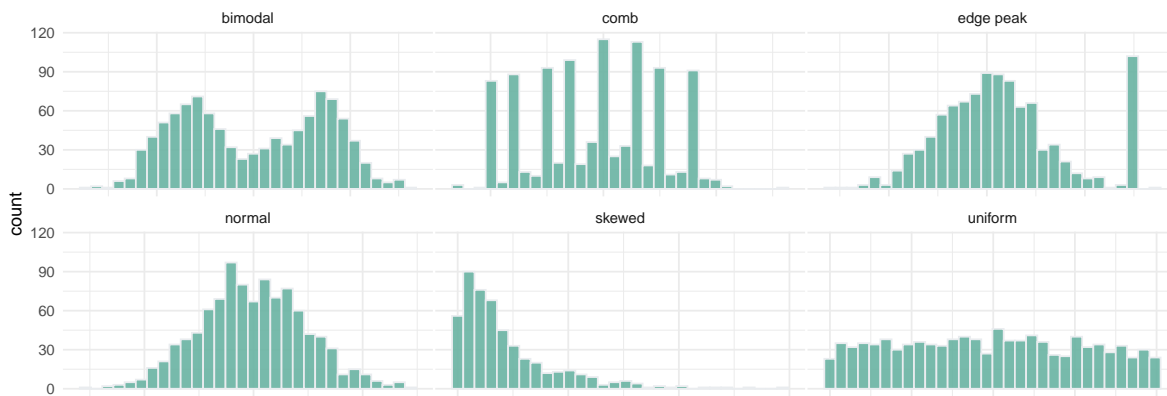
# Build dataset with different distributions
data <- data.frame(
  type = c(rep("edge peak", 1000), rep("comb", 1000), rep("normal", 1000), rep("uniform", 1000)),
  value = c(
    rnorm(900), rep(3, 100), # edge peak
    rnorm(360, sd = 0.5), rep(c(-1, -0.75, -0.5, -0.25, 0, 0.25, 0.5, 0.75), 80), # comb
    rnorm(1000), # normal
    runif(1000), # uniform
  )
)
```

```

    rnorm(500, mean = -2), rnorm(500, mean = 2),      # bimodal
    abs(log(rnorm(1000)))                             # skewed
  )
)

# Represent it
ggplot(data, aes(x = value)) +
  geom_histogram(fill = "#69b3a2", color = "#e9ecef", alpha = 0.9) +
  facet_wrap(~ type, scales = "free_x") +
  theme_minimal() + # Replace hrbrthemes::theme_ipsum() with theme_minimal() from ggplot2
  theme(
    panel.spacing = unit(0.1, "lines"),
    axis.title.x = element_blank(),
    axis.text.x = element_blank(),
    axis.ticks.x = element_blank()
  )

```



Useful for:

Checking this distribution also helps you discovering mistakes in the data. For example, the **comb** distribution can often denote a rounding that has been applied to the variable or another mistake.


As a second step, histogram allow to compare the distribution of **a few** variables. Don't compare more than 3 or 4, it would make the figure cluttered and unreadable. This comparison can be done showing the 2 variables on the same graphic and using transparency.

Example

```
# Libraries
library(ggplot2)

# Build dataset with different distributions
data <- data.frame(
  type = c(rep("variable 1", 1000), rep("variable 2", 1000)),
  value = c(rnorm(1000), rnorm(1000, mean = 4))
)

# Represent it
ggplot(data, aes(x = value, fill = type)) +
  geom_histogram(color = "#e9ecef", alpha = 0.6, position = "identity") +
  scale_fill_manual(values = c("#69b3a2", "#404080")) +
  theme_minimal() + # Replace hrbrthemes::theme_ipsum() with theme_minimal()
  labs(fill = "")
```



Variation

A common variation of the histogram is the mirror histogram: it puts face to face 2 histograms to compare their distribution.


```

# Libraries
library(ggplot2)

# Build dataset
data <- data.frame(
  x = rnorm(1000),
  y = rnorm(1000, mean = 2)
)


# Represent it
ggplot(data) +
  # First histogram for variable 'x'
  geom_histogram(aes(x = x, y = ..density..), binwidth = diff(range(data$x)) / 30, fill = "#69b3a2") +
  geom_label(aes(x = 4.8, y = 0.25, label = "variable1", color = "#69b3a2") +

  # Second histogram for variable 'y' (flipped vertically)
  geom_histogram(aes(x = y, y = -..density..), binwidth = diff(range(data$x)) / 30, fill = "#404080") +
  geom_label(aes(x = 4.8, y = -0.25, label = "variable2", color = "#404080") +

  # Replace hrbrthemes::theme_ipsum() with theme_minimal()
  theme_minimal() +

  # Label the x-axis
  xlab("value of x")

```



Common mistakes

- Try several bin sizes, it can lead to very different conclusions.
- Don't use weird color scheme. It does not give any more insight.
- Don't confound it with a barplot. A barplot gives a value for each group of a categorical variable. Here, we have only a numeric variable and we check its distribution.
- Don't compare more than ~3 groups in the same histogram. The graphic gets cluttered and hardly understandable. Instead use a violin plot, a boxplot, a ridgeline plot or use small multiple.
- Using unequal bin widths

Density Plot

Definition

A density plot is a representation of the distribution of a numeric variable. It uses a kernel density estimate to show the probability density function of the variable. It is a smoothed version of the histogram and is used in the same concept.

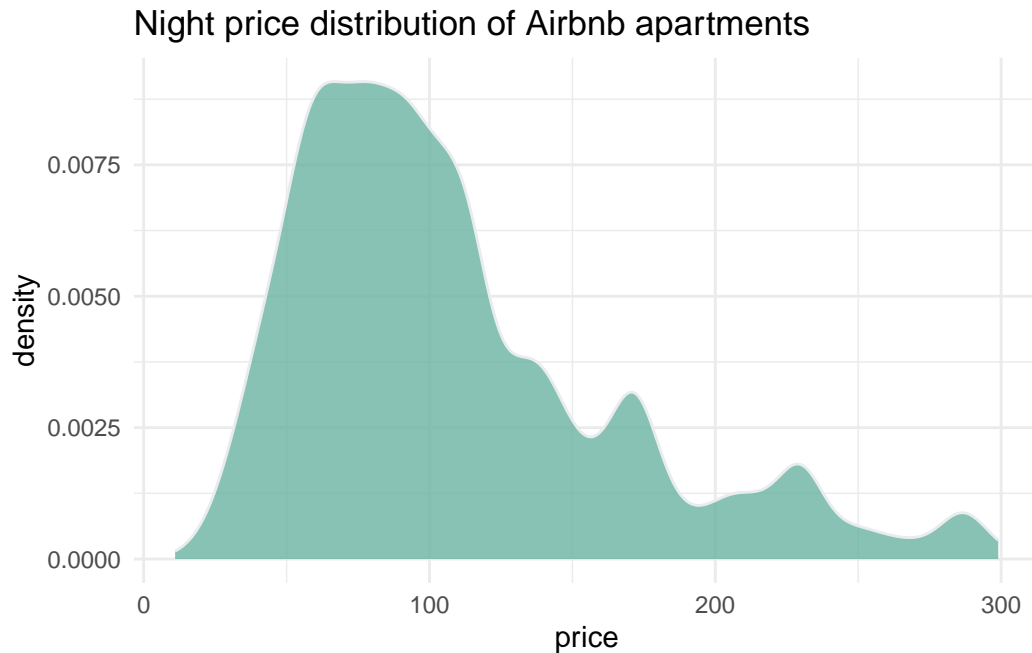
Here is an example showing the distribution of the night price of Rbnb appartements in the south of France:

```
# Libraries
library(ggplot2)

# Load dataset from github
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataa.csv")

# Make the density plot
filtered_data <- subset(data, price < 300)

ggplot(filtered_data, aes(x = price)) +
  geom_density(fill = "#69b3a2", color = "#e9ecef", alpha = 0.8) +
  ggtitle("Night price distribution of Airbnb apartments") +
  theme_minimal()
```



What for

Density plots are used to study the distribution of one or a few variables. Checking the distribution of your variables one by one is probably the first task you should do when you get a new dataset. It delivers a good quantity of information. Several distribution shapes exist, here is an illustration of the 6 most common ones:

Examples

```
# Libraries
library(ggplot2)

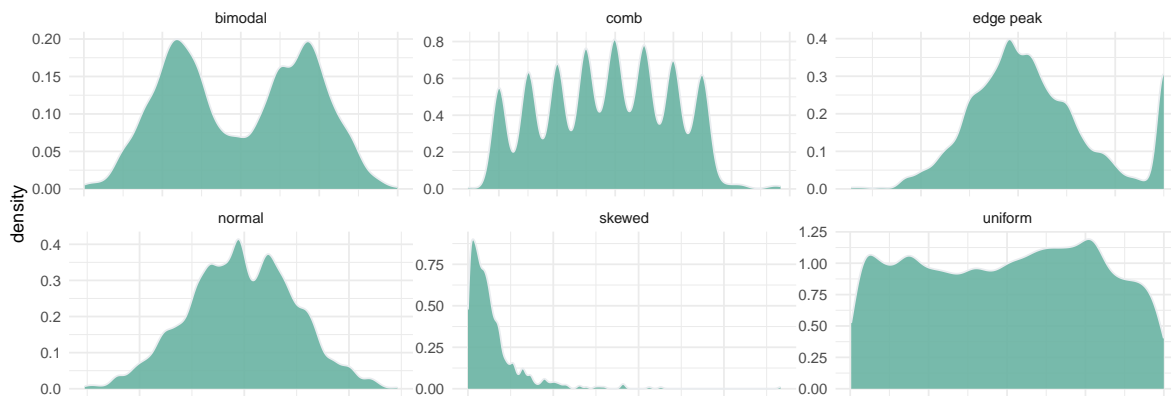
# Build dataset with different distributions
data <- data.frame(
  type = c(rep("edge peak", 1000), rep("comb", 1000), rep("normal", 1000), rep("uniform", 1000)),
  value = c(
    rnorm(900), rep(3, 100), # edge peak
    rnorm(360, sd = 0.5), rep(c(-1, -0.75, -0.5, -0.25, 0, 0.25, 0.5, 0.75), 80), # comb
    rnorm(1000), # normal
    runif(1000), # uniform
  )
)
```

```

    rnorm(500, mean = -2), rnorm(500, mean = 2),      # bimodal
    abs(log(rnorm(1000)))                             # skewed
  )
)

# Represent it
ggplot(data, aes(x = value)) +
  geom_density(fill = "#69b3a2", color = "#e9ecef", alpha = 0.9, adjust = 0.5) +
  facet_wrap(~type, scales = "free") + # Corrected 'scale' to 'scales'
  theme_minimal() + # Replace hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    panel.spacing = unit(0.1, "lines"),
    axis.title.x = element_blank(),
    axis.text.x = element_blank(),
    axis.ticks.x = element_blank()
  )

```



Useful for:

Checking this distribution also helps you discovering mistakes in the data. For example, the comb distribution can often denote a rounding that has been applied to the variable or another mistake.


As a second step, density plots allow to compare the distribution of **a few** variables. Don't compare more than 3 or 4, it would make the figure cluttered and unreadable. This comparison can be done showing the 2 variables on the same graphic and using transparency.

Example

```
# Libraries
library(ggplot2)

# Build dataset with different distributions
data <- data.frame(
  type = c(rep("variable 1", 1000), rep("variable 2", 1000)),
  value = c(rnorm(1000), rnorm(1000, mean = 4))
)

# Represent it
ggplot(data, aes(x = value, fill = type)) +
  geom_density(color = "#e9ecef", alpha = 0.6) +
  scale_fill_manual(values = c("#69b3a2", "#404080")) +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  labs(fill = "")
```



Variation

A common variation of the density plot is the mirror density plot: it puts face to face 2 density plots to compare their distribution.

```

# Libraries
library(ggplot2)

# Build dataset
data <- data.frame(
  x = rnorm(1000),
  y = rnorm(1000, mean = 2)
)


# Represent it
ggplot(data) +
  # First density for variable 'x'
  geom_density(aes(x = x, y = ..density..), fill = "#69b3a2") +
  geom_label(aes(x = 4.5, y = 0.25, label = "variable1"), color = "#69b3a2") +

  # Second density for variable 'y' (flipped vertically)
  geom_density(aes(x = y, y = -..density..), fill = "#404080") +
  geom_label(aes(x = 4.5, y = -0.25, label = "variable2"), color = "#404080") +

  # Replace hrbrthemes::theme_ipsum() with theme_minimal()
  theme_minimal() +

  # Label for the x-axis
  xlab("value of x")

```



Common mistakes

- Play with the bandwidth argument, it can lead to very different conclusions.
- Don't compare more than ~3 groups on the same density plot. The graphic gets cluttered and hardly understandable. Instead use a violin plot, a boxplot, a ridgeline plot or use small multiple.

Scatter Plot

Definition

A scatterplot displays the relationship between 2 numeric variables. For each data point, the value of its first variable is represented on the X axis, the second on the Y axis.

Here is an example considering the price of 1460 apartments and their ground living area. This dataset comes from a Kaggle machine learning competition.

```
# Libraries
library(ggplot2)
library(tidyverse)

# Load dataset from github using base R
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset_1.csv",
                  header = TRUE, sep = ";", c("GrLivArea", "SalePrice"))

# Plot
ggplot(data, aes(x = GrLivArea, y = SalePrice / 1000)) +
  geom_point(color = "#69b3a2", alpha = 0.6) +
  ggtitle("Ground living area partially explains sale price of apartments") +
  theme_minimal() + # Replaced hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    plot.title = element_text(size = 12)
  ) +
  ylab('Sale price (k$)') +
  xlab('Ground living area')
```



What for

A scatterplot is made to study the relationship between 2 variables. Thus it is often accompanied by a correlation coefficient calculation, that usually tries to measure the linear relationship.

However other types of relationship can be detected using scatterplots, and a common task consists to fit a model explaining Y in function of X. Here are a few patterns you can detect doing a scatterplot.

Examples

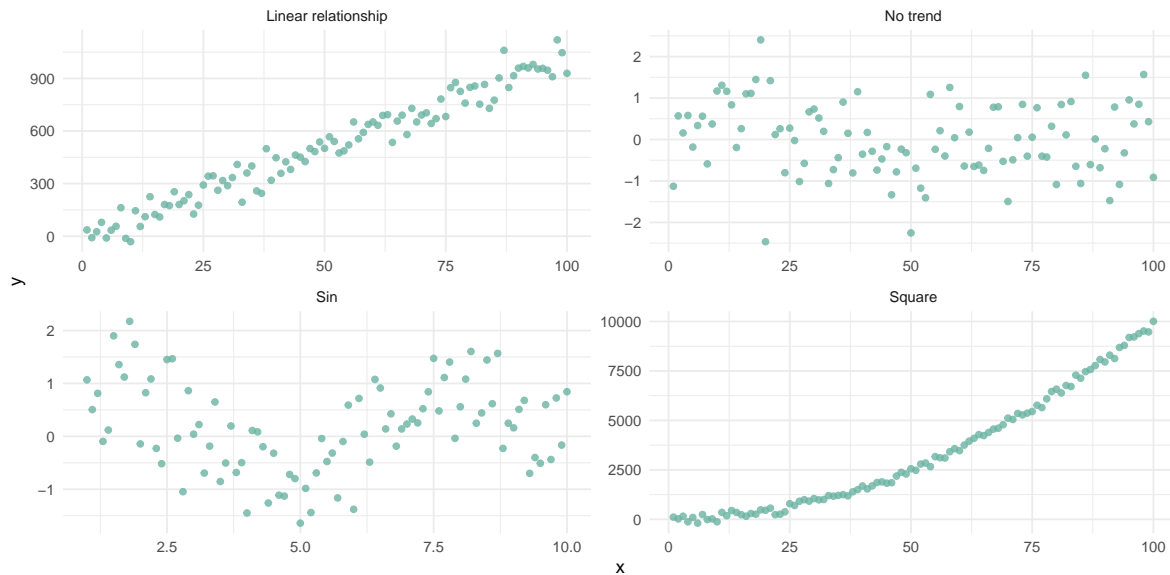
```
# Libraries
library(ggplot2)

# Create data
d1 <- data.frame(x = seq(1, 100), y = rnorm(100), name = "No trend")
d2 <- transform(d1, y = x * 10 + rnorm(100, sd = 60), name = "Linear relationship")
d3 <- transform(d1, y = x^2 + rnorm(100, sd = 140), name = "Square")
```



```
d4 <- data.frame(x = seq(1, 10, 0.1), y = sin(seq(1, 10, 0.1)) + rnorm(91, sd = 0.6), name =
don <- do.call(rbind, list(d1, d2, d3, d4))

# Plot
ggplot(don, aes(x = x, y = y)) +
  geom_point(color = "#69b3a2", alpha = 0.8) +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  facet_wrap(~name, scales = "free") # 'scale' should be 'scales' in facet_wrap()
```



Useful for:

- Detecting relationships between two variables
- Identifying patterns or trends in data
- Spotting outliers or unusual data points
- Comparing different groups or categories within the data

Variation

Scatterplots are sometimes supported by marginal distributions. It indeed adds insight to the graphic, revealing the distribution of both variables:

```

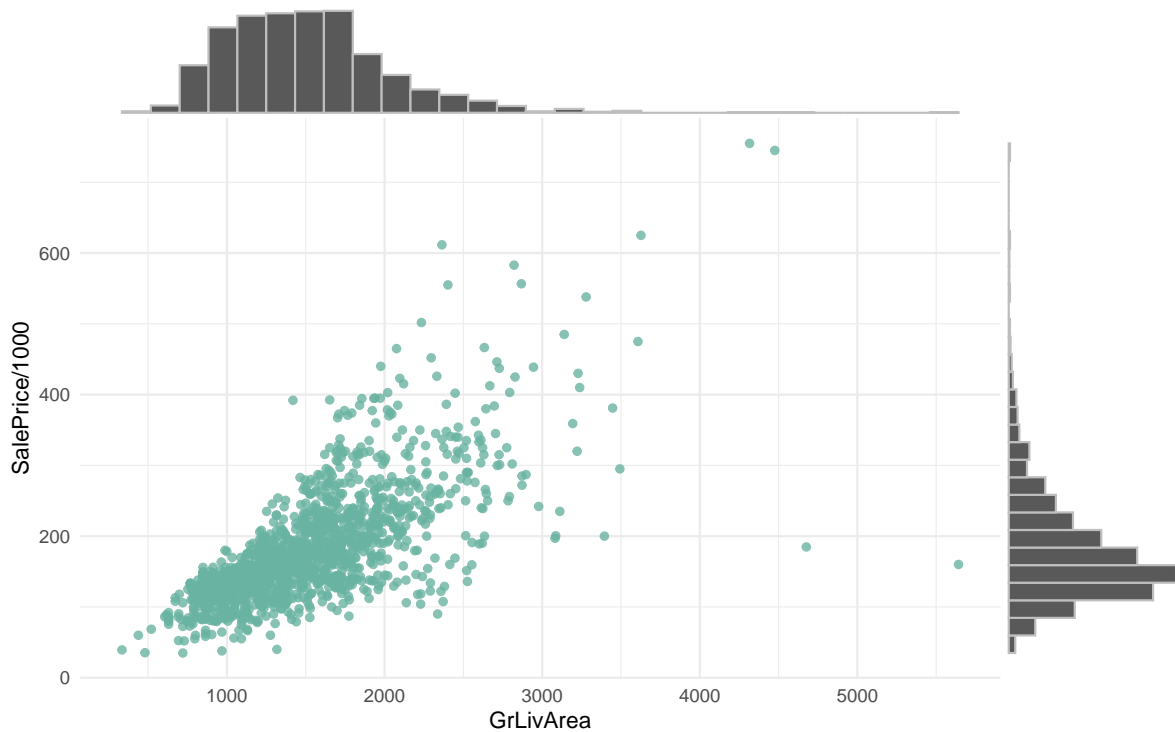
# Libraries
library(ggplot2)
library(ggExtra)
library(tidyverse)

# Load dataset from github using base R
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset_1.csv",
                  header = TRUE, sep = ",")[, c("GrLivArea", "SalePrice")]

# Create a ggplot2 scatterplot
p <- data %>%
  ggplot(aes(x = GrLivArea, y = SalePrice / 1000)) +
  geom_point(color = "#69b3a2", alpha = 0.8) +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    legend.position = "none"
  )

# Add marginal histograms
ggMarginal(p, type = "histogram", color = "grey")

```



Common mistakes

- Overplotting is the most common mistake when sample size is high. There are about 10 different workarounds to fix this issue.
- Don't forget to show subgroups if you have some. Indeed it can reveal important hidden patterns in your data, like in the case of the Simpson paradox.

Histogram

Definition


A barplot (or barchart) is one of the most common types of graphic. It shows the relationship between a numeric and a categoric variable. Each entity of the categoric variable is represented as a bar. The size of the bar represents its numeric value.

Here is an example showing the quantity of weapons exported by the top 20 largest exporters in 2017:

```
# Libraries
library(ggplot2)

# Load dataset from github using base R
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset_1.csv",
                  header = TRUE, sep = ",")

# Barplot
data %>%
  filter(!is.na(Value)) %>%
  arrange(Value) %>%
  tail(20) %>%
  mutate(Country = factor(Country, Country)) %>%
  ggplot(aes(x = Country, y = Value)) +
  geom_bar(stat = "identity", fill = "#69b3a2") +
  coord_flip() +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    panel.grid.minor.y = element_blank(),
    panel.grid.major.y = element_blank(),
    legend.position = "none"
  ) +
  xlab("") +
  ylab("Weapon quantity (SIPRI trend-indicator value)")
```



What for

A barplot shows the relationship between a numeric and a categoric variable. In the previous graphic, each country is a level of the categoric variable, and the quantity of weapon sold is the numeric variable. An ordered barplot is a very good choice here since it displays both the ranking of countries and their specific value.


A barplot can also display values for several levels of grouping. Here's an example of a grouped barplot:

Examples

```
# Libraries
library(ggplot2)
library(babynames)

# Load dataset
data <- babynames %>%
  filter(name %in% c("Ashley", "Amanda", "Jessica", "Patricia", "Linda", "Deborah", "Dorothy"))
  filter(sex == "F")

# A grouped barplot
data %>%
  filter(name %in% c("Ashley", "Patricia", "Betty", "Helen")) %>%
  filter(year %in% c(1920, 1960, 2000)) %>%
  mutate(year = as.factor(year)) %>%
  ggplot(aes(x = year, y = n, fill = name)) +
  geom_bar(stat = "identity", position = "dodge") +
  scale_fill_manual(values = c("#69b3a2", "#404080", "#ffcc00", "#ff5733")) + # Custom colors
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  ylab("Number of babies")
```



Useful for:

- Comparing values across categories

- Showing the distribution of a numeric variable for different groups
- Displaying rankings or ordered data
- Visualizing part-to-whole relationships (in stacked bar charts)

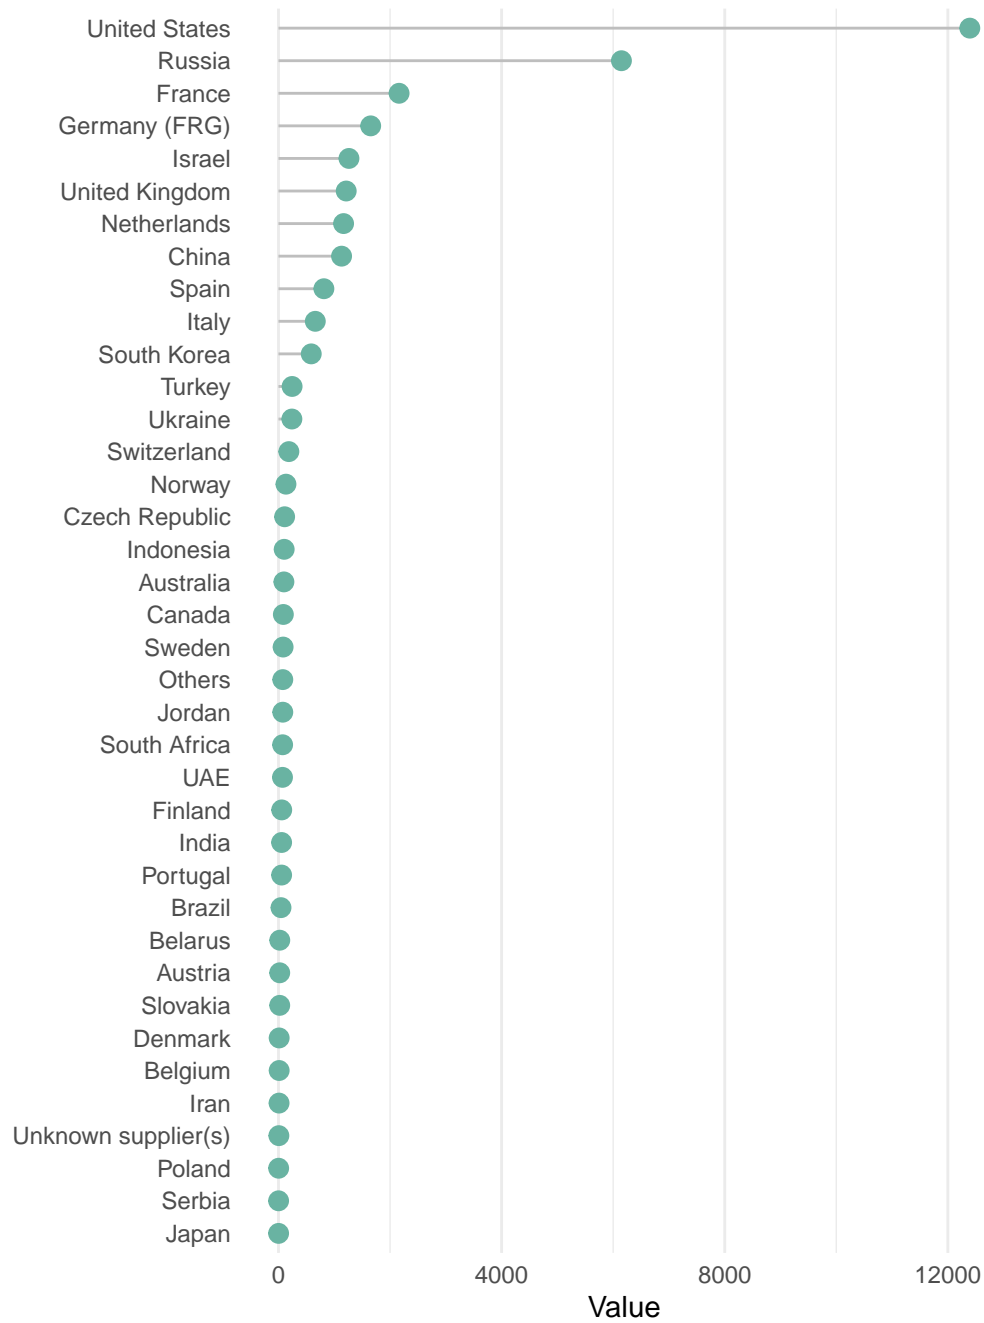
Variation

Lollipop plot: A variation that replaces bars with segments and dots, resulting in a less cluttered figure:

```
# Libraries
library(ggplot2)

# Load dataset from github
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset1.csv",
                  header = TRUE, sep = ",")

# Plot
data %>%
  filter(!is.na(Value)) %>%
  arrange(Value) %>%
  mutate(Country = factor(Country, Country)) %>%
  ggplot(aes(x = Country, y = Value)) +
  geom_segment(aes(x = Country, xend = Country, y = 0, yend = Value), color = "grey") +
  geom_point(size = 3, color = "#69b3a2") +
  coord_flip() +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    panel.grid.minor.y = element_blank(),
    panel.grid.major.y = element_blank(),
    legend.position = "none"
  ) +
  xlab("")
```



Circular barplot: Useful when you have many groups, though it sacrifices some accuracy in comparison:

```

# Libraries
library(ggplot2)
library(dplyr)

# Load dataset from github
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_data")

# Order data
tmp <- data %>%
  filter(!is.na(Value)) %>%
  arrange(desc(Value)) %>%
  mutate(Country = factor(Country, Country))

# Set a number of 'empty bars'
empty_bar = 10

# Add empty bars to the dataset
to_add <- matrix(NA, empty_bar, ncol(tmp))
colnames(to_add) <- colnames(tmp)
tmp <- rbind(tmp, to_add)
tmp$id <- seq(1, nrow(tmp))

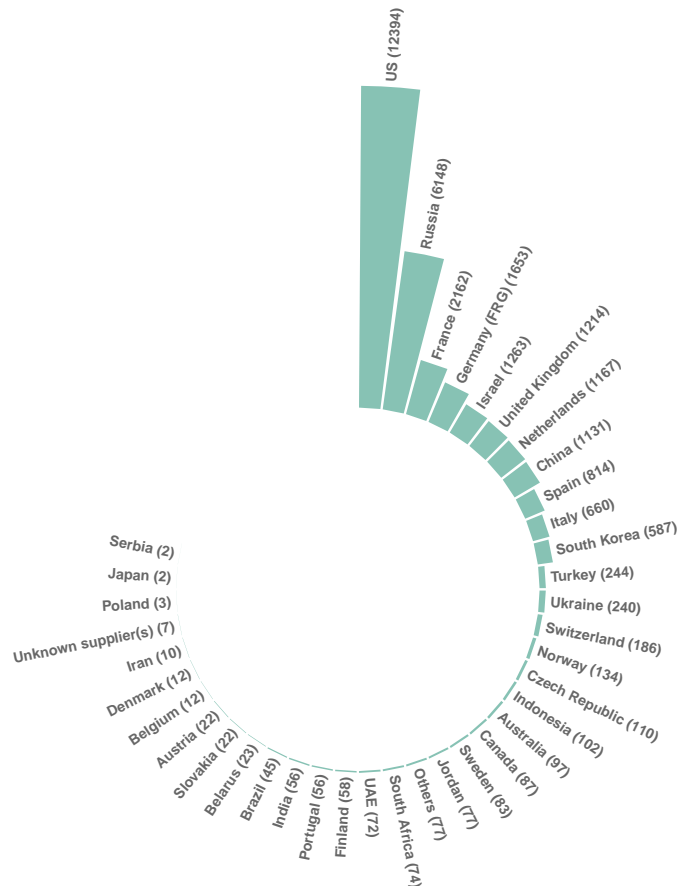
# Get the name and the y position of each label
label_tmp <- tmp
number_of_bar <- nrow(label_tmp)
angle <- 90 - 360 * (label_tmp$id - 0.5) / number_of_bar # Adjust angles for label position
label_tmp$hjust <- ifelse(angle < -90, 1, 0)
label_tmp$angle <- ifelse(angle < -90, angle + 180, angle)
label_tmp$Country <- gsub("United States", "US", label_tmp$Country)
label_tmp$Country <- paste(label_tmp$Country, " (", label_tmp$Value, ")", sep = "")

# Make the plot
ggplot(tmp, aes(x = as.factor(id), y = Value)) + # x is treated as a factor
  geom_bar(stat = "identity", fill = alpha("#69b3a2", 0.8)) +
  ylim(-7000, 13000) +
  theme_minimal() +
  theme(
    axis.text = element_blank(),
    axis.title = element_blank(),
    panel.grid = element_blank(),
    plot.margin = unit(rep(-1, 4), "cm")
  ) +

```



```
coord_polar(start = 0) + # Polar coordinates for a circular plot
geom_text(data = label_tmp, aes(x = id, y = Value + 200, label = Country),
          color = "black", fontface = "bold", alpha = 0.6, size = 2.5,
          angle = label_tmp$angle, hjust = label_tmp$hjust, inherit.aes = FALSE) +
geom_text(aes(x = 24, y = 8000, label = "Who sells more weapons?"),
          color = "black", inherit.aes = FALSE, data = data.frame())
```



Who sells more weapons?

Common mistakes

- Do not confound barchart with histogram. A histogram has only a numeric variable as input and shows its distribution.
- Order your bars. If the levels of your categoric variable have no obvious order, order the bars following their values.
- Several values per group? Don't use a barplot. Even with error bars, it hides information and other types of graphics like boxplot or violin are much more appropriate.

Line Chart

Definition

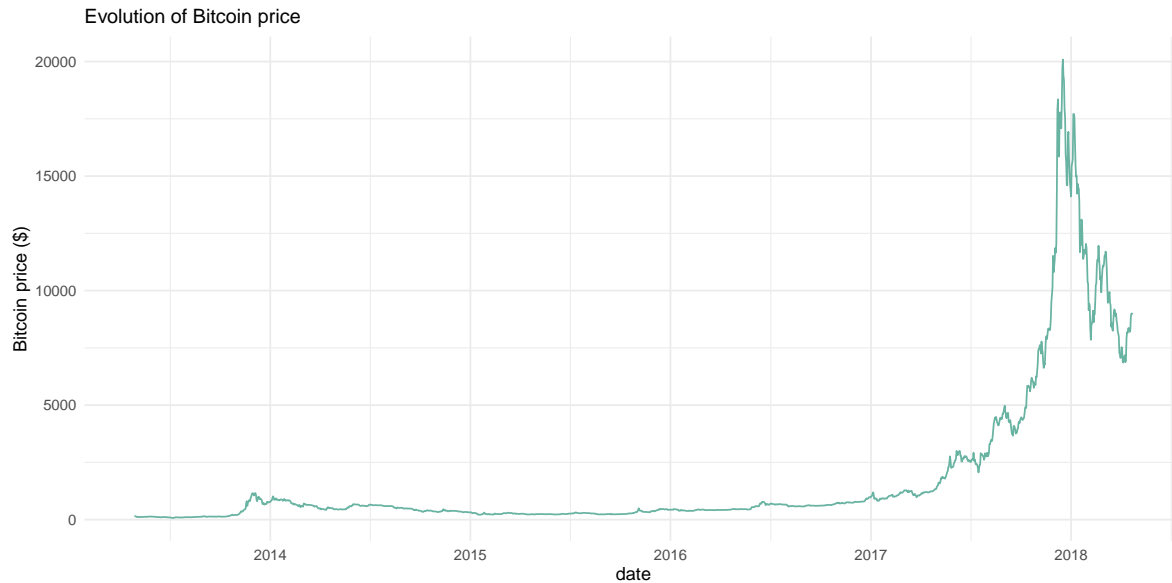
A line chart or line graph displays the evolution of one or several numeric variables. Data points are connected by straight line segments. It is similar to a scatter plot except that the measurement points are ordered (typically by their x-axis value) and joined with straight line segments. A line chart is often used to visualize a trend in data over intervals of time – a time series – thus the line is often drawn chronologically.

Here is an example showing the evolution of the bitcoin price between April 2013 and April 2018:

```
# Libraries
library(ggplot2)

# Load dataset from github
data <- read.table("https://raw.githubusercontent.com/holtzy/data_to_viz/master/Example_dataset1.csv",
                  header = TRUE)
data$date <- as.Date(data$date)

# Plot
ggplot(data, aes(x = date, y = value)) +
  geom_line(color = "#69b3a2") +
  ggtitle("Evolution of Bitcoin price") +
  ylab("Bitcoin price ($)") +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    plot.title = element_text(size = 12)
  )
```



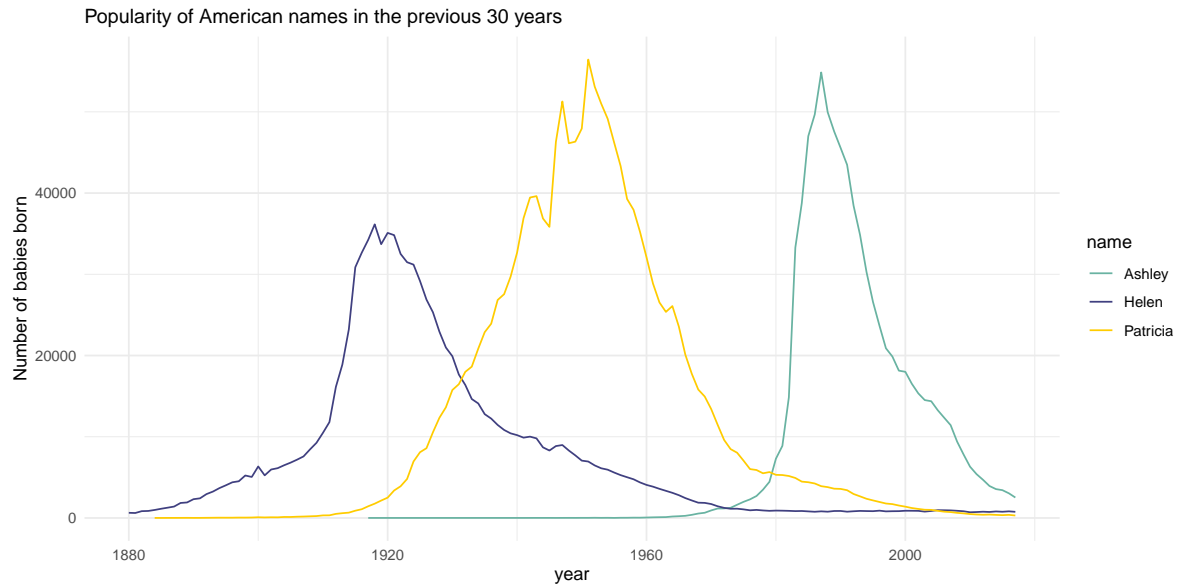
What for

Line charts can be used to show the evolution of one or several variables. Here is an example showing the evolution of three baby name frequencies in the US between 1880 and 2015:

```
# Libraries
library(ggplot2)
library(babynames)

# Load dataset from babynames
don <- babynames %>%
  filter(name %in% c("Ashley", "Patricia", "Helen")) %>%
  filter(sex == "F")

# Plot
ggplot(don, aes(x = year, y = n, group = name, color = name)) +
  geom_line() +
  scale_color_manual(values = c("#69b3a2", "#404080", "#ffcc00")) + # Custom colors to repl
  ggtitle("Popularity of American names in the previous 30 years") +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  ylab("Number of babies born") +
  theme(
    plot.title = element_text(size = 12)
  )
```



Useful for:

- Visualizing trends over time
- Comparing multiple variables or categories over a continuous axis
- Showing the rate of change between data points
- Identifying patterns, cycles, or anomalies in data


Variation

If the number of data points is low, it is advised to represent each individual observation with a dot:

```
# Libraries
library(ggplot2)

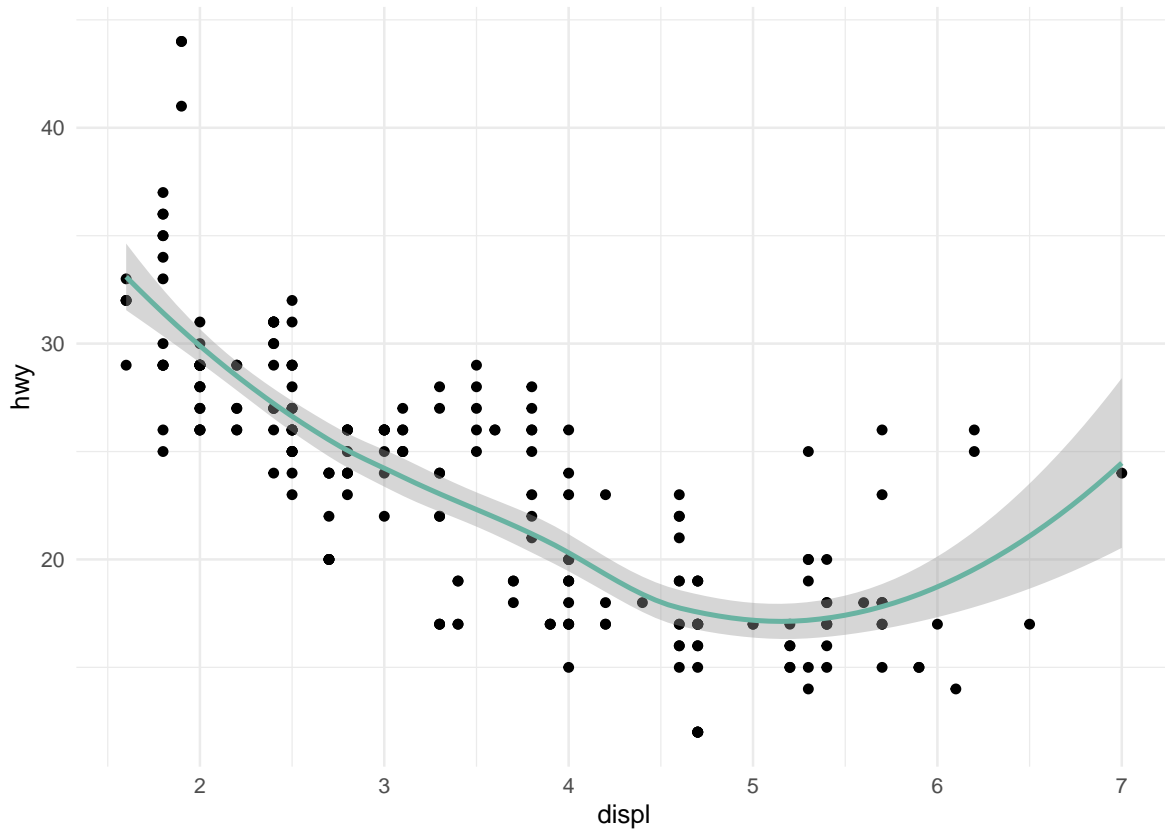
# Assuming `data` is already loaded and contains `date` and `value`
# Using tail to get the last 10 entries for the plot
data %>%
  tail(10) %>%
  ggplot(aes(x = date, y = value)) +
  geom_line(color = "#69b3a2") +
  geom_point(color = "#69b3a2", size = 4) +
  ggtitle("Line chart with points") +
```

```
ylab("Bitcoin price ($)") +  
theme_minimal() # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
```



Lines are also used to show trends in a scatterplot. Here is an example using Smoothed conditional means and showing confidence interval around it:

```
# Libraries  
library(ggplot2)  
  
# Plot using the mpg dataset  
ggplot(mpg, aes(displ, hwy)) +  
  geom_point() +  
  geom_smooth(color = "#69b3a2") +  
  theme_minimal()
```



Common caveats

The line chart is subject to a lot of discussion concerning potential caveats:

- To cut or not to cut the Y axis? Whether or not the Y axis must start at 0 is a hot topic leading to intense debates. Generally, line plots do not need to start at 0 since it allows observing patterns more efficiently.

Line Chart

Common caveats

The line chart is subject to a lot of discussion concerning potential caveats:

- To cut or not to cut the Y axis? Whether or not the Y axis must start at 0 is a hot topic leading to intense debates. Generally, line plots do not need to start at 0 since it allows observing patterns more efficiently.

```

# Libraries
library(ggplot2)
library(patchwork)

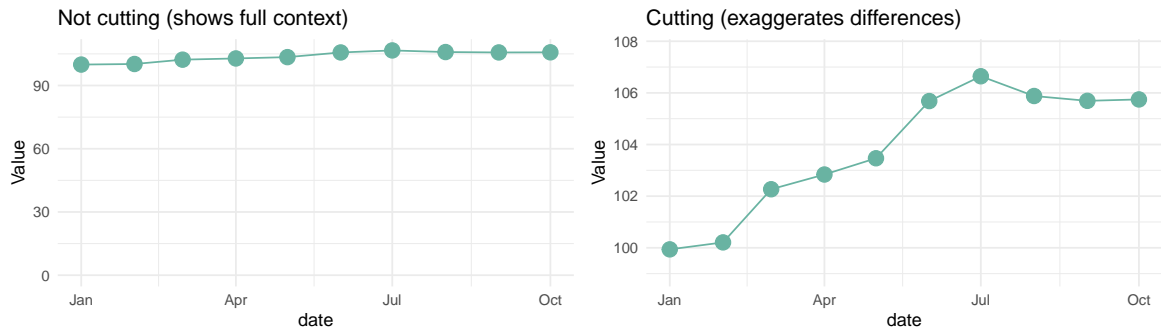
# Create a small dataset with less dramatic changes
set.seed(123)
small_data <- data.frame(
  date = seq(as.Date("2023-01-01"), by = "month", length.out = 10),
  value = 100 + cumsum(rnorm(10, mean = 0.5, sd = 1))
)

# Plot 1: Full context, no cutting
p1 <- small_data %>%
  ggplot(aes(x = date, y = value)) +
  geom_line(color = "#69b3a2") +
  geom_point(color = "#69b3a2", size = 4) +
  ggtitle("Not cutting (shows full context)") +
  ylab("Value") +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  ylim(0, max(small_data$value))

# Plot 2: Cutting y-axis to exaggerate differences
p2 <- small_data %>%
  ggplot(aes(x = date, y = value)) +
  geom_line(color = "#69b3a2") +
  geom_point(color = "#69b3a2", size = 4) +
  ggtitle("Cutting (exaggerates differences)") +
  ylab("Value") +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  ylim(min(small_data$value) - 1, max(small_data$value) + 1)

# Combine plots using patchwork
p1 + p2

```



- If you need to compare the evolution of 2 different variables, avoid using dual axis. Dual axis can show very different results depending on what range you apply to the axis.
- Mind the spaghetti chart: too many lines make the chart unreadable.
- Think about the aspect ratio of the graphic, extreme ratios make the chart unreadable.

Word Cloud

Definition

A word cloud (also called tag cloud or weighted list) is a visual representation of text data. Words are usually single words, and the importance of each is shown with font size or color.

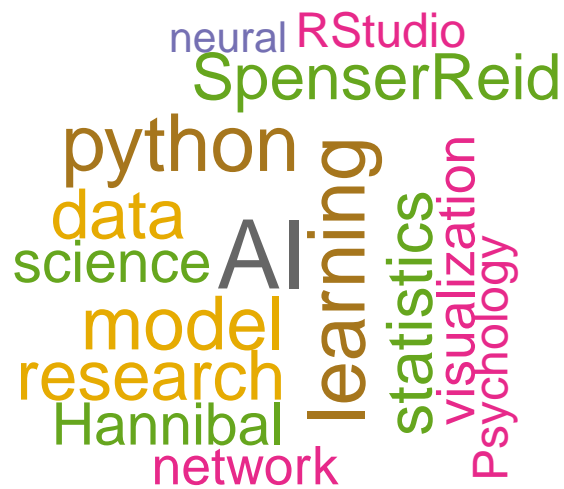
Here is an example showing the most frequent words used by Nekfeu, a famous French rapper, in a few of his songs:

```
# Libraries
library(wordcloud)

# Example data (random words with frequencies)
words <- c("data", "science", "AI", "Psychology", "learning",
           "neural", "network", "SpenserReid", "model", "Hannibal",
           "research", "statistics", "RStudio", "python", "visualization")

frequencies <- c(10, 8, 15, 6, 12, 5, 7, 9, 11, 8, 10, 9, 6, 12, 7)

# Create a word cloud
set.seed(123) # For reproducibility
wordcloud(words = words, freq = frequencies, min.freq = 1,
           scale = c(3, 0.5), colors = brewer.pal(8, "Dark2"),
           random.order = FALSE, rot.per = 0.35)
```

What for

Word clouds are useful for:

- Quickly perceiving the most prominent terms
- Locating a term alphabetically to determine its relative prominence
- Creating visually appealing representations of text data

They are widely used in media and well understood by the public.

Useful for:

- Summarizing large amounts of text data
- Highlighting key themes or topics in a dataset
- Creating engaging visuals for presentations or reports
- Comparing word frequencies across different texts or sources

Example

Here's an alternative representation using a lollipop plot, which addresses some of the limitations of word clouds:

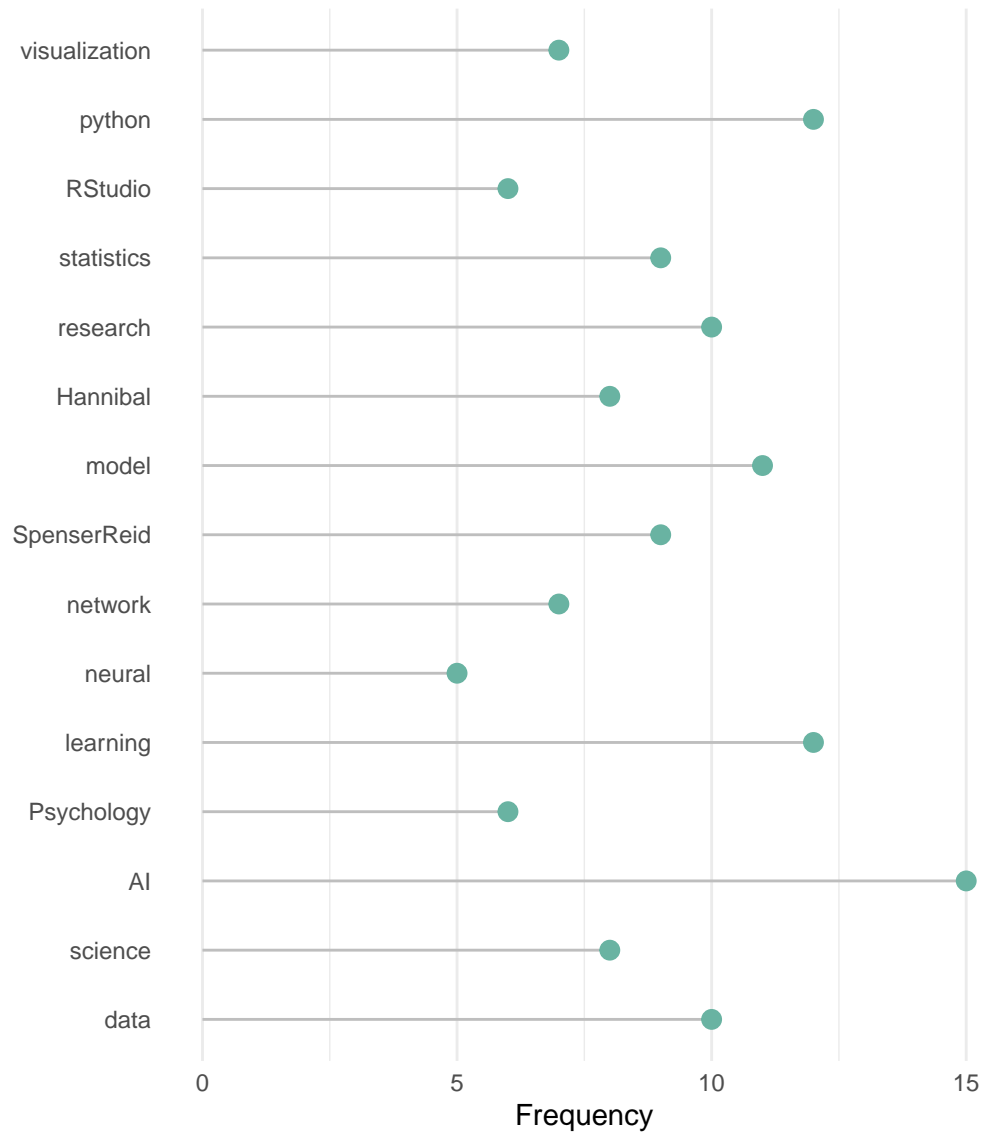
```
# Load necessary libraries
library(dplyr)
library(ggplot2)

# Example data (random words with frequencies)
words <- c("data", "science", "AI", "Psychology", "learning",
           "neural", "network", "SpenserReid", "model", "Hannibal",
           "research", "statistics", "RStudio", "python", "visualization")

frequencies <- c(10, 8, 15, 6, 12, 5, 7, 9, 11, 8, 10, 9, 6, 12, 7)

# Create a data frame from the provided words and frequencies
word_data <- data.frame(word = words, freq = frequencies)

# Plot using ggplot2
word_data %>%
  mutate(word = factor(word, word)) %>%
  ggplot(aes(x = word, y = freq)) +
  geom_segment(aes(x = word, xend = word, y = 0, yend = freq), color = "grey") +
  geom_point(size = 3, color = "#69b3a2") +
  coord_flip() +
  theme_minimal() + # Replacing hrbrthemes::theme_ipsum() with theme_minimal()
  theme(
    panel.grid.minor.y = element_blank(),
    panel.grid.major.y = element_blank(),
    legend.position = "none"
  ) +
  xlab("") +
  ylab("Frequency")
```



Variation

Many variations exist for word clouds:

- Different shapes, sometimes using the shape of an object related to the topic
- Varying text orientation, font, size, and colors
- Interactive word clouds that change or provide additional information on hover

Common mistakes

- Relying too heavily on word clouds for accurate data representation
- Ignoring the limitations of area as a metaphor for numeric values
- Not accounting for the bias created by longer words appearing larger
- Using word clouds when more precise visualizations (like bar charts or lollipop plots) would be more appropriate
- Creating overly complex or cluttered word clouds that are difficult to read or interpret

Word clouds, while visually appealing, are often criticized for their lack of accuracy in conveying information. Consider using them primarily for aesthetic purposes or in conjunction with more precise data visualization methods.

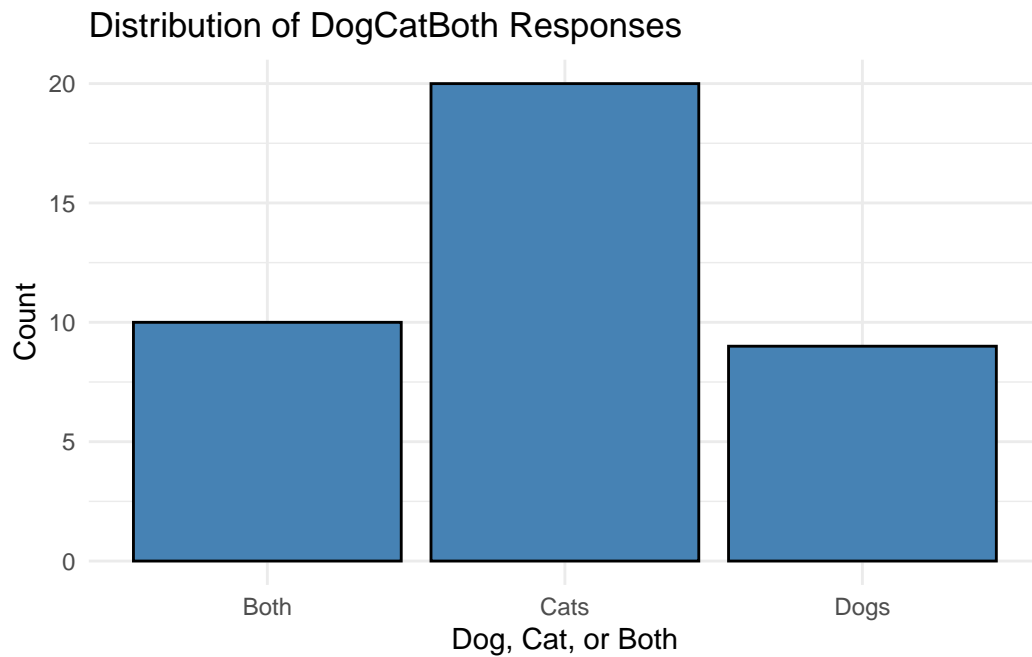
Psychology Data from Year 1

```
# Load necessary libraries
library(ggplot2)

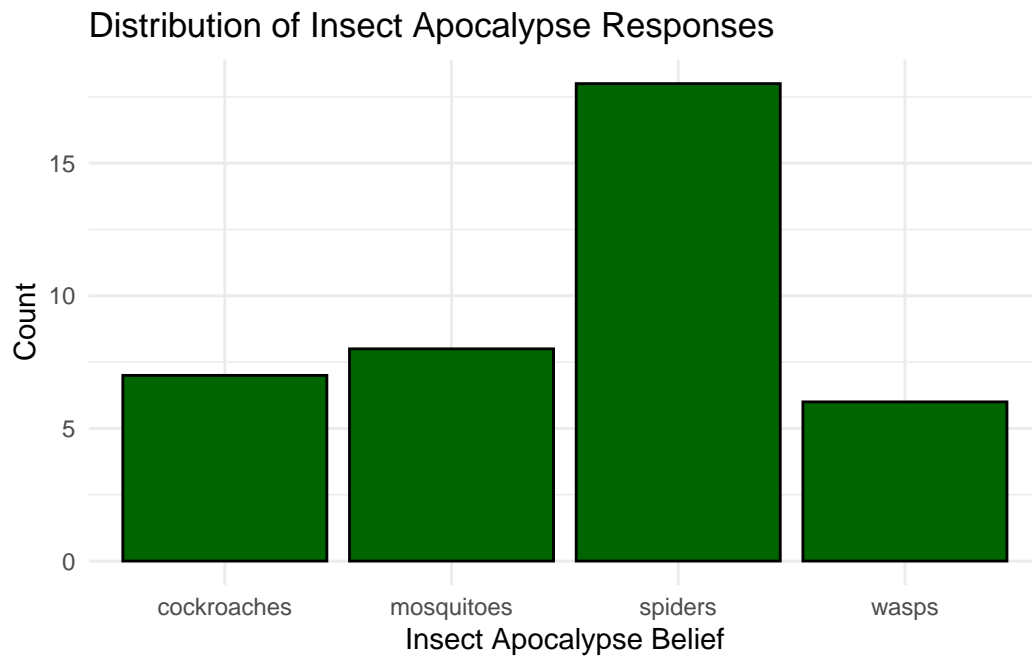
# Load the dataset
data <- read.csv("materials/data/Y1W3_data.csv")

# Remove all NA values
data <- na.omit(data)

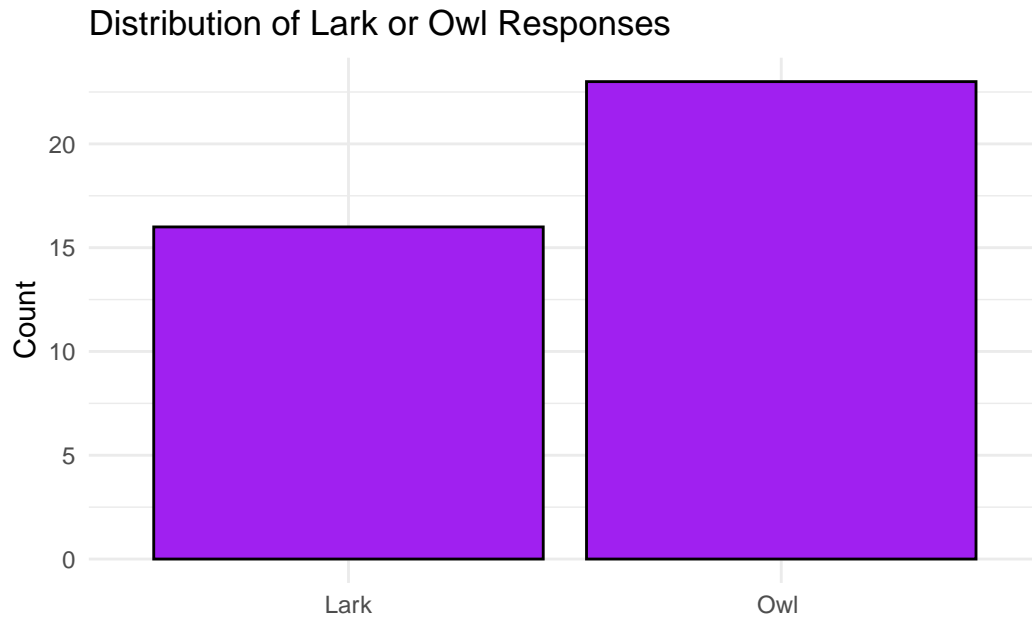
# Create a bar chart for the 'DogCatBoth' variable
ggplot(data, aes(x = DogCatBoth)) +
  geom_bar(fill = "steelblue", color = "black") +
  theme_minimal() +
  labs(title = "Distribution of DogCatBoth Responses",
       x = "Dog, Cat, or Both",
       y = "Count")
```



```
# Create a bar chart for the 'InsectApocalypse' variable
ggplot(data, aes(x = InsectApocalypse)) +
  geom_bar(fill = "darkgreen", color = "black") +
  theme_minimal() +
  labs(title = "Distribution of Insect Apocalypse Responses",
       x = "Insect Apocalypse Belief",
       y = "Count")
```



```
# Create a bar chart for the 'LarkorOwl' variable
ggplot(data, aes(x = LarkorOwl)) +
  geom_bar(fill = "purple", color = "black") +
  theme_minimal() +
  scale_x_discrete(labels = function(x) substr(x, 1, 4)) +
  labs(title = "Distribution of Lark or Owl Responses",
       x = "",
       y = "Count")
```



```
# Count for each column  
cat("Count for DogCatBoth: ", sum(!is.na(data$DogCatBoth)), "\n")
```

Count for DogCatBoth: 39

```
cat("Count for InsectApocalypse: ", sum(!is.na(data$InsectApocalypse)), "\n")
```

Count for InsectApocalypse: 39

```
cat("Count for LarkorOwl: ", sum(!is.na(data$LarkorOwl)), "\n")
```

Count for LarkorOwl: 39

```
# Load required libraries  
library(ggplot2)  
library(dplyr)  
library(tidyr)  
library(scales)
```

Attaching package: 'scales'

The following object is masked from 'package:purrr':

discard

The following object is masked from 'package:readr':

col_factor

```
# Read the data
data <- read.csv("materials/data/Y1W3_data.csv")

# Select relevant columns and remove rows with NA
tipi_data <- data %>%
  select(LarkorOwl, E_TIPI, A_TIPI, C_TIPI, ES_TIPI, O_TIPI) %>%
  na.omit()

# Shorten labels to Lark and Owl
tipi_data$LarkorOwl <- ifelse(grepl("Lark", tipi_data$LarkorOwl), "Lark", "Owl")

# Reshape the data for radar plot
radar_data <- tipi_data %>%
  gather(key = "Trait", value = "Score", -LarkorOwl) %>%
  mutate(Trait = factor(Trait, levels = c("E_TIPI", "A_TIPI", "C_TIPI", "ES_TIPI", "O_TIPI")))

# Calculate mean scores for each trait by Lark/Owl preference
mean_scores <- radar_data %>%
  group_by(LarkorOwl, Trait) %>%
  summarize(Mean_Score = mean(Score)) %>%
  ungroup()
```

`summarise()` has grouped output by 'LarkorOwl'. You can override using the
`.groups` argument.

```
# Create a function to generate coordinates for the radar chart
coord_radar <- function(theta = "x", start = 0, direction = 1) {
  theta <- match.arg(theta, c("x", "y"))
  r <- if (theta == "x") "y" else "x"
  ggproto("CordRadar", CoordPolar, theta = theta, r = r,
    start = start, direction = sign(direction),
    is_linear = function(coord) TRUE)
}
```



```

# Create the radar plot
radar_plot <- ggplot(mean_scores, aes(x = Trait, y = Mean_Score, group = LarkorOwl, color = LarkorOwl)) +
  geom_polygon(aes(fill = LarkorOwl), alpha = 0.25) +
  geom_line(size = 1) +
  geom_point(size = 3) +
  scale_y_continuous(limits = c(0, 5), breaks = seq(0, 5, 1)) +
  scale_color_manual(values = c("Lark" = "#FF6B6B", "Owl" = "#4ECDC4")) +
  scale_fill_manual(values = c("Lark" = "#FF6B6B", "Owl" = "#4ECDC4")) +
  coord_radar() +
  facet_wrap(~ LarkorOwl, ncol = 2) +
  theme_minimal() +
  theme(
    axis.text.x = element_text(size = 10),
    axis.text.y = element_blank(),
    axis.title = element_blank(),
    panel.grid.major = element_line(color = "gray", size = 0.3),
    panel.grid.minor = element_blank(),
    legend.position = "none"
  ) +
  labs(title = "TIPI Trait Scores by Lark/Owl Preference",
       subtitle = "E: Extraversion, A: Agreeableness, C: Conscientiousness, ES: Emotional Stability")

```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.

Warning: The `size` argument of `element_line()` is deprecated as of ggplot2 3.4.0.
 i Please use the `linewidth` argument instead.

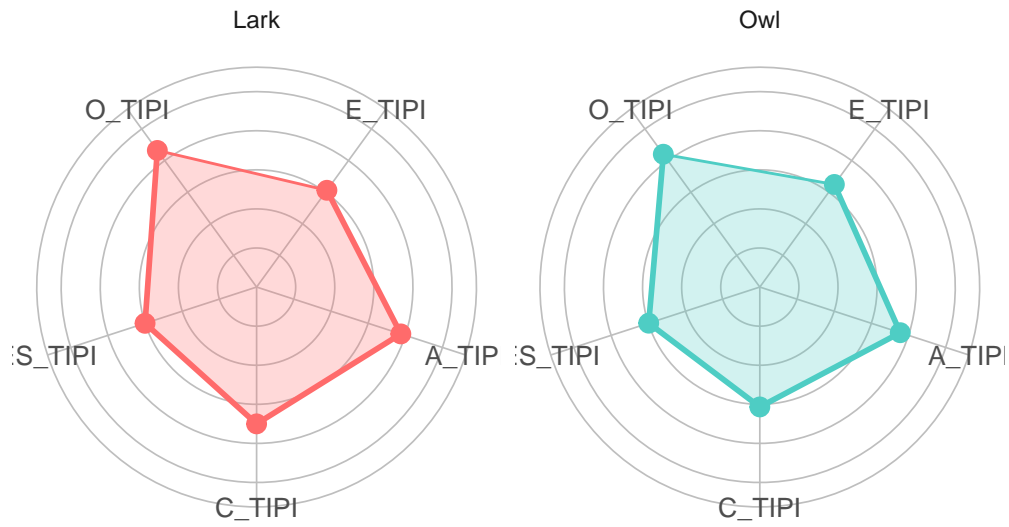
```

# Display the plot
print(radar_plot)

```

TIPi Trait Scores by Lark/Owl Preference

E: Extraversion, A: Agreeableness, C: Conscientiousness, ES: Emotional Stabi



```
# Save the plot as a PNG file
ggsave("tipi_radar_plots.png", radar_plot, width = 12, height = 6, dpi = 300)
```